

# EECE 5639 Computer Vision I

---

Lecture 5

**Filtering**

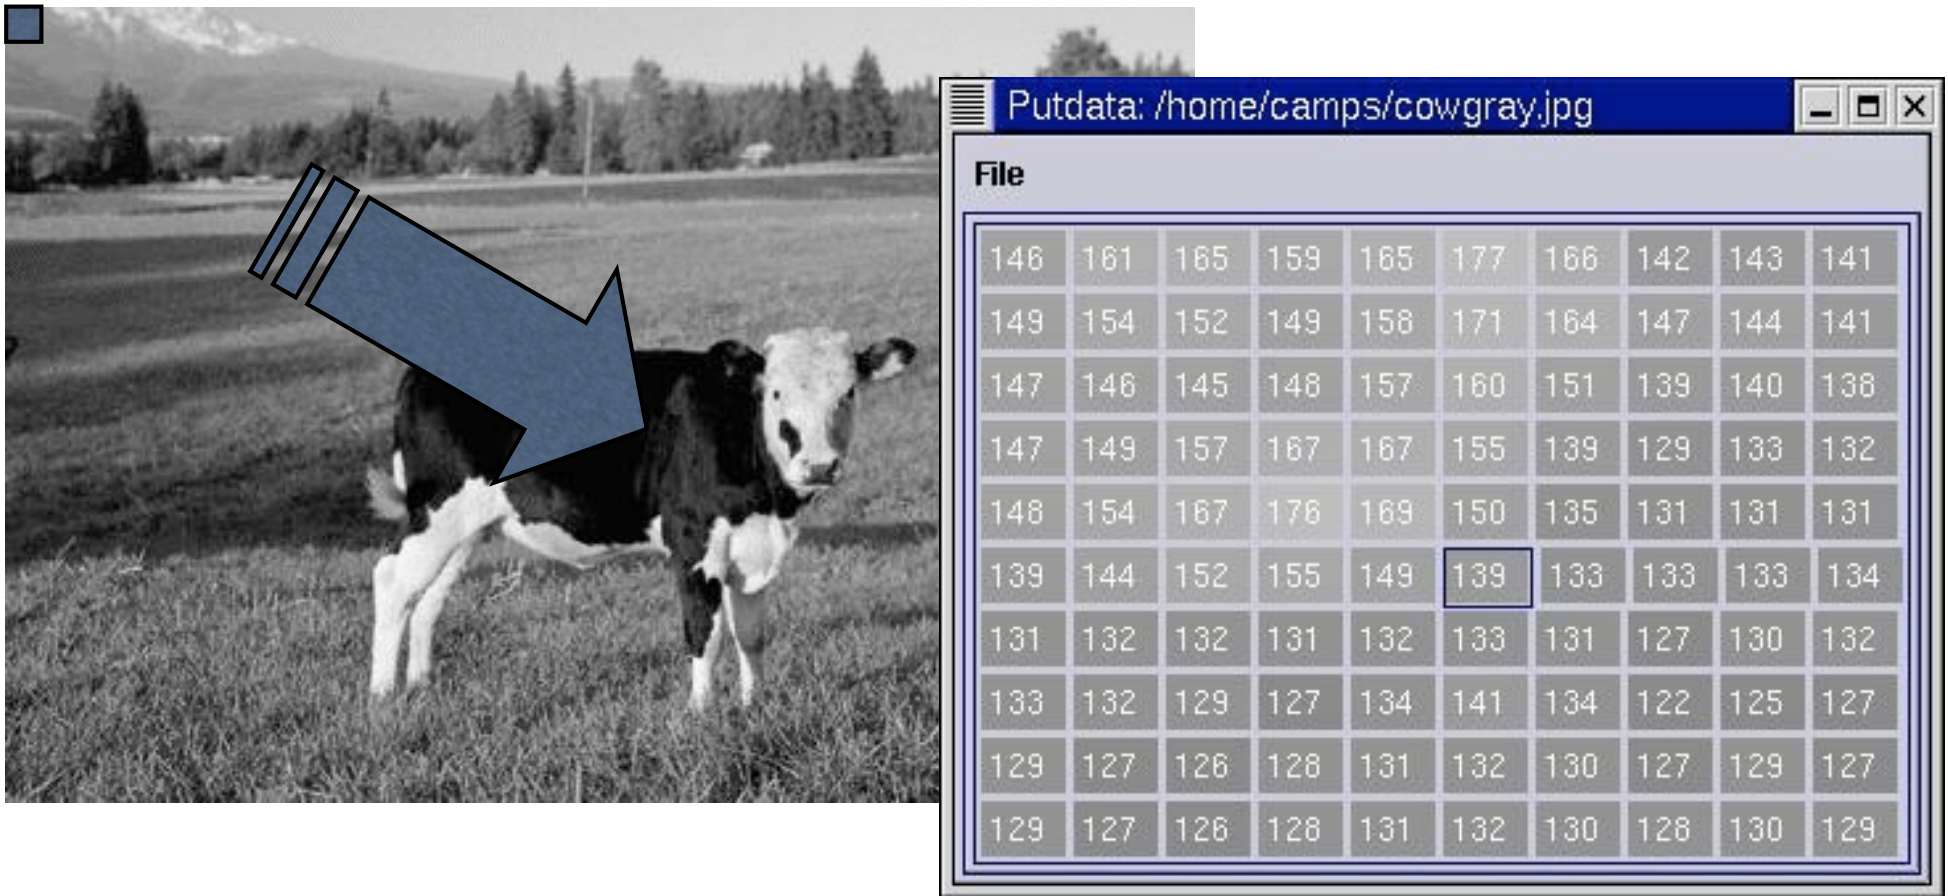
Next Class

**Edges, Corners**

# Image processing: Filtering

# Digital Images

are 2D arrays (matrices) of numbers:




# Digital Images


Array of numbers (pixels)

Typically integers 0-255 (unsigned byte)

Pixels have neighbors




4-neighbors



8-neighbors

# Image Paths

A path is a sequence of pixel indices  $(i_0, j_0)(i_1, j_1) \dots (i_n, j_n)$  such that  $(i_k, j_k)$  is a neighbor of  $(i_{k+1}, j_{k+1})$



4-connected  
path

8-connected path

# Levels of Computation

---

## Point level

Output based only on a single point

Ex.: thresholding

## Local level

Output based on a neighborhood

Ex. : smoothing and edge detection

## Global level

Output based on the whole image

Ex.: Fourier transform and histogram

## Object level

Output based on pixels that belong to an object

# Spatial Filtering

# Spatial Filtering

---

- Use of spatial masks (kernels, filters, templates, windows) for image processing (spatial filters)
- Linear and nonlinear filters
- Spatial Filters include:
  - Sharpening
  - Smoothing
  - Edge detection
  - Noise removal
  - etc



# Linear Filters

- General process:
  - Form new image whose pixels are a weighted sum of original pixel values, using the same set of weights at each point.
- Properties
  - Output is a linear function of the input
  - Output is a shift-invariant function of the input (i.e. shift the input image two pixels to the left, the output is shifted two pixels to the left)
- Example: smoothing by averaging
  - form the average of pixels in a neighborhood
- Example: smoothing with a Gaussian
  - form a weighted average of pixels in a neighborhood
- Example: finding a derivative
  - form a weighted average of pixels in a neighborhood

**Note: The “Linear” in “Linear Filters” means linear combination of neighboring pixel values.**

# Image Filtering

---

**Low-pass** filters eliminate or attenuate high frequency components in the frequency domain (sharp image details), and result in image blurring.

**High-pass** filters attenuate or eliminate low-frequency components (resulting in sharpening edges and other sharp details).

**Band-pass** filters remove selected frequency regions between low and high frequencies (for image restoration, not enhancement).


# Spatial Filtering

---

Operations are performed directly on the pixels in the spatial domain.

The process involves sweeping a mask on the image and performing at each point a set of predefined operations on the pixels overlapped by the mask.

# Basics of Spatial Filtering



**FIGURE 3.32** The mechanics of spatial filtering. The magnified drawing shows a  $3 \times 3$  mask and the image section directly under it; the image section is shown displaced out from under the mask for ease of readability.

Ex of a 3x3 mask


# Linear Spatial Filtering: CORRELATION

Linear filtering of an  $M \times N$  image  $f(x,y)$  with a filter  $w(s,t)$  of size  $m \times n$  is given by:

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

$$a = (m - 1)/2; \quad b = (n - 1)/2$$

Where  $m$  and  $n$  are odd numbers



This is similar to convolution ... and it is often referred as “convolving with a mask”

# Linear Spatial Filtering: CORRELATION

Alternative notation:

**FIGURE 3.33**

Another representation of a general  $3 \times 3$  spatial filter mask.

$w_1$	$w_2$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$

$$R = \sum_{i=1}^{m \times n} w_i z_i$$

Where  $z_i$  are the values of the input image under the mask

# Correlation Example

---

1	1	1
-1	2	1
-1	-1	1

h

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

f

# Correlation Example

*Step 1*

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

h

0	0	0		
0	4	2	2	3
0	-2	1	3	3
	2	2	1	2
	1	3	2	2

f



5			

$f*h$



## Step 2

h

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

0	0	0	
-2	4	2	3
-2	-1	3	3
2	2	1	2
1	3	2	2

f



5	4		

f\*h

# Step 3

h

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

	0	0	0
2	-2	4	3
2	-1	-3	3
2	2	1	2
1	3	2	2

f



5	4	4	

f\*h

# Step 4

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

h

		0	0	0
2	2	-2	6	0
2	1	-3	-3	0
2	2	1	2	
1	3	2	2	



5	4	4	-2

f

f\*h

# Step 5

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

h

0	2	2	2	3
0	4	1	3	3
0	-2	2	1	2
	1	3	2	2

f



5	4	4	-2
9			

f\*h

# Step 6

h

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

2	2	2	3
-2	2	3	3
-2	-2	1	2
1	3	2	2

f




5	4	4	-2
9	6		

f\*h

And so on ...


# Practical Issue: Border Handling

- Border issues:
  - When applying convolution with a  $K \times K$  kernel, the result is undefined for pixels closer than  $K$  pixels from the border of the image




- Options:


Warp around



Expand/Pad



Crop



- Reflection at border also a useful option!

# Correlation vs Convolution



Correlation:

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

↓                      ↓

$$a = (m - 1)/2; \quad b = (n - 1)/2$$

Convolution:

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t)$$

↓                      ↓

FLIP FIRST!

$$a = (m - 1)/2; \quad b = (n - 1)/2$$



# Correlation vs Convolution

---



**If the mask is symmetric, then there is no difference.**

# Correlation and Convolution in MATLAB

---



Could use `conv` and `conv2`, but newer versions use:

`Imfilter(image,template{,option1,option2,...})`

Boundary options: constant, symmetric, replicate, circular

Output size options: same as image, or full size (includes partial values computed when mask is off the image).

Corr or conv option: convolution rotates the template (as we have discussed, correlation does not).

Type “help imfilter” on command line for more details

# Smoothing Filters

# Image Noise

---

Images are noisy

Noise is anything in the image that we are not interested in


Examples:

- Fluctuations of pixel values


- Numerical errors

- Clutter

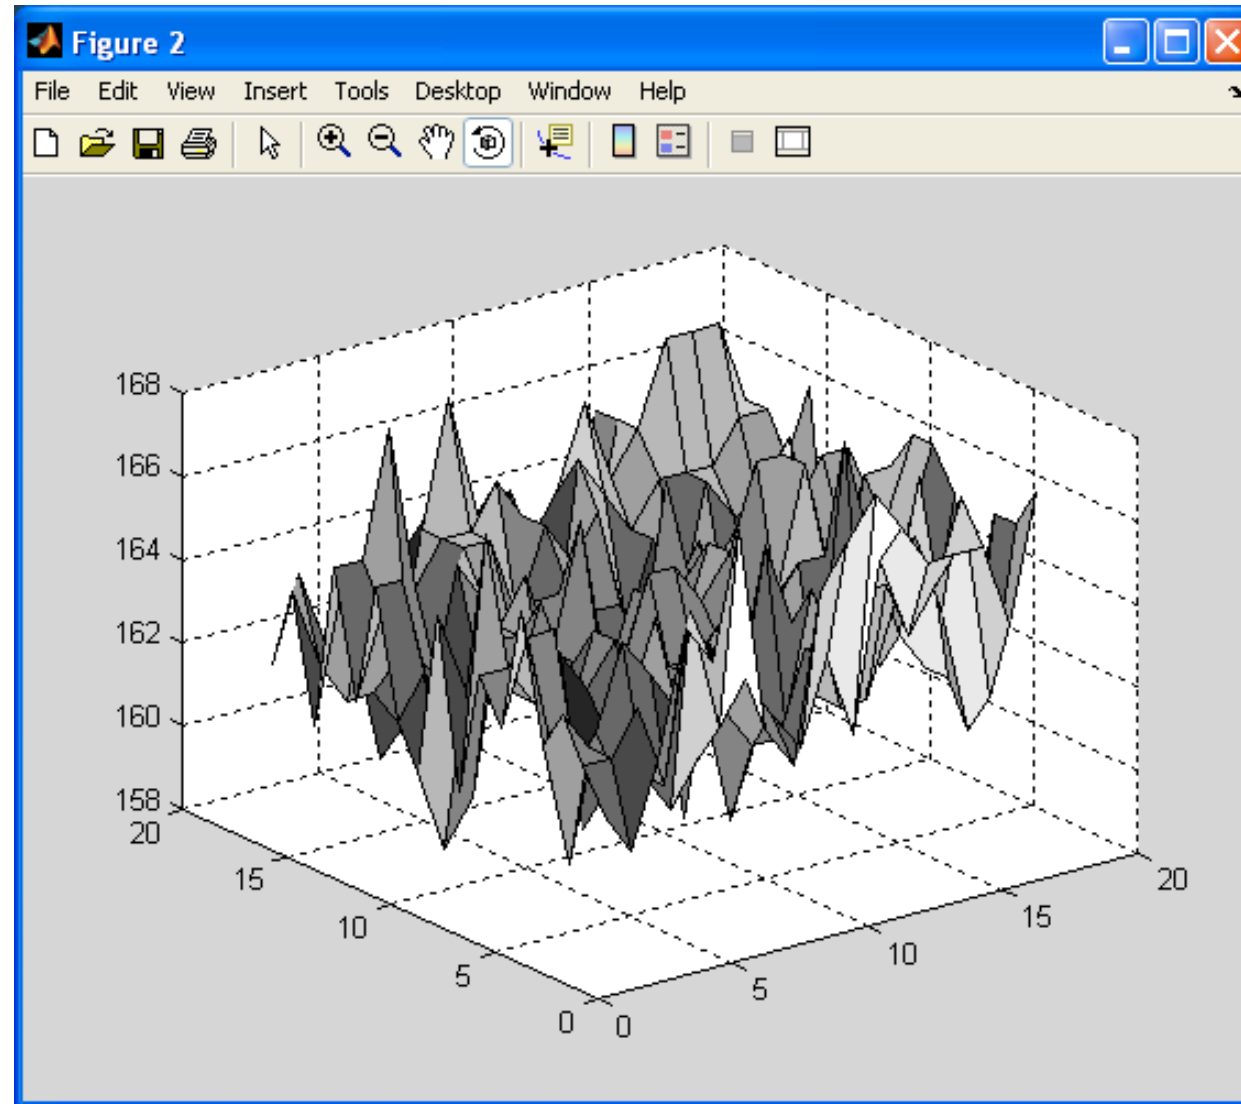
# Images as Surfaces



Surface height  
proportional to  
pixel grey value  
(dark=low, light=high)



# Examples



Mean = 164    Std = 1.8

# Where does noise come from?

---

Light fluctuations

Sensor noise

Quantization effects

Finite precision

# Modeling Noise




---

We are interested in RANDOM noise.

Deterministic noise (ex: hardware defects) can be corrected.



# Dealing with Noise



We want something more like this!

# Probability Review

# Intuitive Development

---

Intuitively, the probability of an event **a** could be defined as:

$$P(a) = \lim_{n \rightarrow \infty} \frac{N(a)}{n}$$

Where  $N(a)$  is the number that event **a** happens in  $n$  trials

# More Formal:

---

$\Omega$  is the **Sample Space**:

Contains all possible outcomes of an experiment

$\omega$  in  $\Omega$  is a single outcome

$A$  in  $\Omega$  is a set of outcomes of interest

1.  $P(A) \geq 0 \forall A \in \Omega$

2.  $P(\Omega) = 1$

3.  $A_i \cap A_j = \emptyset \forall i, j \Rightarrow P(\cup_{i=1}^n A_i) = \sum_{i=1}^n P(A_i)$

4.  $P(\emptyset) = 0$

# Independence

---


The probability of independent events A, B and C is given by:

$$P(ABC) = P(A)P(B)P(C)$$

A and B are **independent**, if knowing that A has happened does not say anything about B happening

# Conditional Probability

One of the most useful concepts!



$$P(A|B) = \frac{P(AB)}{P(B)}$$

# Bayes Theorem

---


Provides a way to convert a-priori probabilities to a-posteriori probabilities:

$$P(A|B) = \frac{P(AB)}{P(B)}$$

$$P(A|B)P(B) = P(B|A)P(A)$$

# Using Partitions:

If events  $A_i$  are mutually exclusive and partition  $\Omega$



$$A_i \cap A_j = \emptyset \forall i, j$$

$$\bigcup_{i=1, n} A_i = \Omega$$


$$P(B) = \sum_{i=1}^n P(A_i \cap B)$$



# Random Variables

---

A (scalar) random variable  $X$  is a function that maps the outcome of a random event into real scalar values



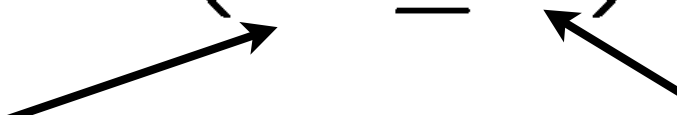
# Random Variables Distributions

---

## Cumulative Probability Distribution (CDF):

$$F_X(x) = P(X \leq x)$$

Random Variable      value



- Probability Density Function (PDF):

$$p_X(x) = \frac{dF_X(x)}{dx}$$

# Random Distributions:

---

From the two previous equations:

$$\int_{-\infty}^{\infty} p_X(x) dx = 1.0$$

**“The area under the curve is equal to 1.0.”**

# Statistical Characterizations

---

Expectation (Mean Value, First Moment):

$$E(X) = \int_{-\infty}^{\infty} xp_X(x)dx$$

“weighted average”

- Second Moment:

$$E(X^2) = \int_{-\infty}^{\infty} x^2 p_X(x)dx$$

# Statistical Characterizations

## Variance of X:

“how far is x from the mean”

$$\begin{aligned} \text{Var}(X) &= E\{[X - E(X)]^2\} \\ &= \int_{-\infty}^{\infty} (x - E[X])^2 p_X(x) dx \\ &= \underbrace{E[X^2]}_{\text{Second Moment}} - \underbrace{(E[X])^2}_{(\text{First Moment})^2} \end{aligned}$$

- Standard Deviation of X:

$$\sigma_X = \sqrt{\text{Var}(X)}$$

# Mean Estimation from Samples

---

Given a set of  $N$  samples from a distribution, we can estimate the mean of the distribution by:

$$\mu = \frac{1}{N} \sum_{i=1}^N x_i$$

# Variance Estimation from Samples

---

Given a set of  $N$  samples from a distribution, we can estimate the variance of the distribution by:


$$\sigma^2 = \frac{1}{N-1} \sum_{i=1}^N (x_i - \mu)^2$$

# Uniform Distribution

---

A R.V.  $X$  that is uniformly distributed between  $x_1$  and  $x_2$  has density function:

$$p_X(x) = \begin{cases} \frac{1}{x_2 - x_1} & x_1 \leq x \leq x_2 \\ 0 & \text{otherwise} \end{cases}$$






# Gaussian (Normal) Distribution

---




A R.V.  $X$  that is normally distributed has density function:

$$p_X(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp -\frac{(x - \mu)^2}{2\sigma^2}$$



... back to images

# Dealing with Noise



We want something  
more like this!

# Image Noise Models

---

Additive noise:

Most commonly used

$$I(i, j) = \hat{I}(i, j) + N(i, j)$$

Multiplicative noise:

$$I(i, j) = \hat{I}(i, j) \times N(i, j)$$

Impulsive noise (salt and pepper):


$$I(i, j) = \begin{cases} \hat{I}(i, j) & \text{if } x < l \\ i_{\min} + y(i_{\max} - i_{\min}) & x \geq l \end{cases}$$

# Additive Noise Models


---

## Gaussian

Usually, zero-mean, uncorrelated



## Uniform



# Measuring Noise

---

Noise Amount:  $\text{SNR} = \sigma_s / \sigma_n$

Noise Estimation:

Given a sequence of images  $I_0, I_1, \dots, I_{N-1}$

$$\bar{I}(i, j) = \frac{1}{N} \sum_{k=0}^{N-1} I_k(i, j)$$

$$\sigma(i, j) = \sqrt{\frac{1}{N-1} \sum_{k=0}^{N-1} (\bar{I}(i, j) - I_k(i, j))^2}$$

$$\sigma_n = \frac{1}{RC} \sum_{i=0}^{R-1} \sum_{j=0}^{C-1} \sigma(i, j)$$

# How can we reduce noise?

---

Image acquisition noise due to light fluctuations and sensor noise can be reduced by acquiring a sequence of images and averaging them.

WHY?

# Smoothing Filters



# Smoothing Spatial Filters

---

They are used for blurring and noise reduction.

Blurring is performed as pre-processing to remove small detail or bridge curve gaps

Noise reduction can be done by linear or nonlinear filtering

# Linear Smoothing Filters

---

They are simply averaging filters: they compute the average of the filters under the mask.

They reduce sharp transitions.

They are low pass filters.

# Average Filter

Mask with **positive** entries, that **sum 1**.

Replaces each pixel with an average of its neighborhood.

If all weights are equal, it is called a BOX filter.

F

1/9

1	1	1
1	1	1
1	1	1

# Example:

**I**

10	11	10	0	0	1
9	10	11	1	0	1
10	9	10	0	2	1
11	10	9	10	9	11
9	10	11	9	99	11
10	9	9	11	10	10

**F**


1	1	1
1	1	1
1	1	1

**O**

X	X	X	X	X	X
X	10				X
X					X
X					X
X					X
X	X	X	X	X	X


$$1/9.(10 \times 1 + 11 \times 1 + 10 \times 1 + 9 \times 1 + 10 \times 1 + 11 \times 1 + 10 \times 1 + 9 \times 1 + 10 \times 1) = 1/9.(90) = 10$$

# Example:




$$1/9.(10 \times 1 + 0 \times 1 + 0 \times 1 + 11 \times 1 + 1 \times 1 + 0 \times 1 + 10 \times 1 + 0 \times 1 + 2 \times 1) = 1/9.(34) = 3.7778$$

# Example:



$$1/9.(10 \times 1 + 9 \times 1 + 11 \times 1 + 9 \times 1 + 99 \times 1 + 11 \times 1 + 11 \times 1 + 10 \times 1 + 10 \times 1) = 1/9.(180) = 20$$

# Example:



$$1/9.(10 \times 1 + 0 \times 1 + 2 \times 1 + 9 \times 1 + 10 \times 1 + 9 \times 1 + 11 \times 1 + 9 \times 1 + 99 \times 1) = 1/9.(159) = 17.6667$$

# Does it reduce noise?

---

Intuitively, takes out small variations.

Consider the Image pixel values under the mask and the corresponding output at the center:

$$I_i = \hat{I}_i + N_i \quad i = 1, \dots, mn$$

$$O = \frac{1}{mn} \sum_i^{mn} (\hat{I}_i + N_i)$$

Is the output better? How?



# Does it reduce noise?

---

Assume that the noise in the image is uncorrelated, zero mean, with stdev sigma.

The **expected value** of a pixel before filtering is:

$$E[I_i] = E[\hat{I}_i + N_i] = E[\hat{I}_i] + E[N_i] = \hat{I}_i + 0 = \hat{I}_i$$
$$i = 1, \dots, mn$$

# Does it reduce noise?

Assume that the noise in the image is uncorrelated, zero mean, with stdev sigma.

The **expected value** of a pixel after filtering is:

$$\begin{aligned} E[O] &= E\left[\frac{1}{mn} \sum_i^{mn} (\hat{I}_i + N_i)\right] \\ &= \frac{1}{mn} \sum_i^{mn} E[\hat{I}_i] + \frac{1}{mn} \sum_i^{mn} E[N_i] \\ &= \frac{1}{mn} \sum_i^{mn} \hat{I}_i + 0 \\ &= \frac{1}{mn} \sum_i^{mn} \hat{I}_i \end{aligned}$$

# Does it reduce noise?

---

Assume that the noise in the image is uncorrelated, zero mean, with stdev sigma.

The **variance** of a pixel before filtering is:

$$E[(I_i - E[I_i])^2] = E[(\hat{I}_i + N_i - \hat{I}_i)^2] = E[N_i^2] = \sigma^2$$
$$i = 1, \dots, mn$$

# Does it reduce noise?

Assume that the noise in the image is uncorrelated, zero mean, with stdev sigma.

The **variance** of the pixel after filtering is:

$$\begin{aligned} E[(O - E[O])^2] &= E\left[\left(\frac{1}{mn} \sum_i^{mn} (\hat{I}_i + N_i) - \frac{1}{mn} \sum_i^{mn} (\hat{I}_i)\right)^2\right] \\ &= \frac{1}{(mn)^2} E\left[\left(\sum_i^{mn} N_i\right)^2\right] \\ &= \frac{1}{(mn)^2} mn \sigma^2 = \frac{\sigma^2}{mn} \end{aligned}$$

# How big should the mask be?

---

The bigger the mask,  
more neighbors contribute.  
smaller noise variance of the output.  
bigger noise spread.  
more blurring.  
more expensive to compute.

# Weighted Average Filter

Gives more weight at the central pixel and less weights to the neighbors.

The farther away the neighbors, the smaller the weight.  
Less blurring of edges


$\frac{1}{9} \times$	1	1	1
	1	1	1
	1	1	1

$\frac{1}{16} \times$	1	2	1
	2	4	2
	1	2	1

$$g(x, y) = \frac{\sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)}{\sum_{s=-a}^a \sum_{t=-b}^b w(s, t)}$$

# Gaussian Filter

A particular case of weighted averaging:  
The coefficients are a 2D Gaussian.



$$w(s, t) = K e^{-\frac{s^2 + t^2}{2\sigma^2}}$$

(0,0) is the center of the mask

$\sigma$  determines how fast the weights decay

K is s.t. the sum of the coefficients is 1

# How big should the mask be?

---

The std. dev of the Gaussian  $\sigma$  determines the amount of smoothing.

The samples should adequately represent a Gaussian

For a 98.76% of the area, we need

$$m = 5\sigma$$

$$5.(1/\sigma) \leq 2\pi \Rightarrow \sigma \geq 0.796, m \geq 5$$



# Efficient Implementation

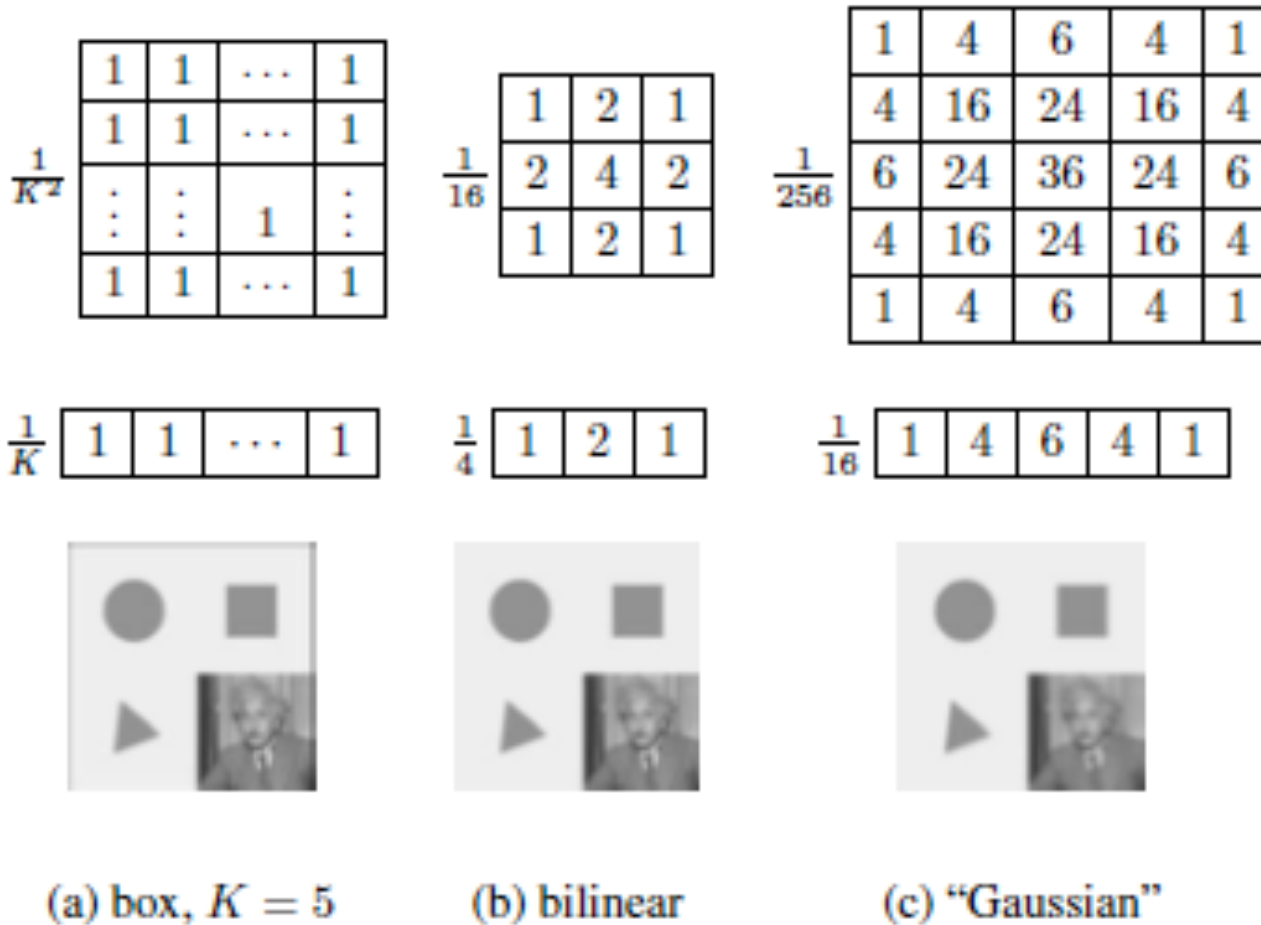
---

Both, the BOX filter and the Gaussian filter are separable:


First convolve each row with a 1D filter

Then convolve each column with a 1D filter.


# Separable Filters



# Effects of increasing mask size



**FIGURE 3.35** (a) Original image, of size  $500 \times 500$  pixels. (b)–(f) Results of smoothing with square averaging filter masks of sizes  $n = 3, 5, 9, 15$ , and  $35$ , respectively. The black squares at the top are of sizes 3, 5, 9, 15, 25, 35, 45, and 55 pixels, respectively; their borders are 25 pixels apart. The letters at the bottom range in size from 10 to 24 points, in increments of 2 points; the large letter at the top is 60 points. The vertical bars are 5 pixels wide and 100 pixels high; their separation is 20 pixels. The diameter of the circles is 25 pixels, and their borders are 15 pixels apart; their gray levels range from 0% to 100% black in increments of 20%. The background of the image is 10% black. The noisy rectangles are of size  $50 \times 120$  pixels.



a b c

**FIGURE 3.36** (a) Image from the Hubble Space Telescope. (b) Image processed by a  $15 \times 15$  averaging mask. (c) Result of thresholding (b). (Original image courtesy of NASA.)

---

# Efficient Implementation: Integral Image

If an image is going to be repeatedly convolved with different box filters a pre-computed summed area table can save computations for future use.

$$s(i, j) = \sum_{k=0}^i \sum_{l=0}^j f(k, l)$$

3	2	7	2	3
1	5	1	3	4
5	1	3	5	1
4	3	2	1	6
2	4	1	4	8

# Efficient Implementation: Integral Image

If an image is going to be repeatedly convolved with different box filters a pre-computed summed area table can save computations for future use.

$$s(i, j) = \sum_{k=0}^i \sum_{l=0}^j f(k, l)$$

3	2	7	2	3
1	5	1	3	4
5	1	3	5	1
4	3	2	1	6
2	4	1	4	8

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

$$s(i, j) = s(i - 1, j) + s(i, j - 1) - s(i - 1, j - 1) + f(k, l)$$

# Efficient Implementation: Integral Image

If an image is going to be repeatedly convolved with different box filters a pre-computed summed area table can save computations for future use.

$$s(i, j) = \sum_{k=0}^i \sum_{l=0}^j f(k, l)$$

3	2	7	2	3
1	5	1	3	4
5	1	3	5	1
4	3	2	1	6
2	4	1	4	8

(a)  $S = 24$

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

(b)  $s = 28$

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

(c)  $S = 24$

$$S = 48 - 14 - 13 + 3$$

# Limitations of averaging

---

Signal frequencies shared with noise are lost, resulting in blurring.

Impulsive noise is diffused but not removed.

It spreads pixel values, resulting in blurring.



# Non-linear Filtering

---

Replace each pixel with the MEDIAN value of all the pixels in the neighborhood.

# Example:

I

10	11	10	0	0	1
9	10	11	1	0	1
10	9	10	0	2	1
11	10	9	10	9	11
9	10	11	9	99	11
10	9	9	11	10	10

O

X	X	X	X	X	X
X	10				X
X					X
X					X
X					X
X	X	X	X	X	X

10, 11, 10, 9, 10, 11, 10, 9,  
10

sort  
→

9, 9, 10, 10, 10, 10, 10, 11, 11

median

# Example:

I

10	11	10	0	0	1
9	10	11	1	0	1
10	9	10	0	2	1
11	10	9	10	9	11
9	10	11	9	99	11
10	9	9	11	10	10

O

X	X	X	X	X	X
X	10	10	1	1	X
X					X
X					X
X					X
X	X	X	X	X	X

11, 10, 0, 10, 11, 1, 9, 10  
0

sort  
→

0, 0, 1, 9, 10, 10, 10, 11, 11

median

# Example:

I

10	11	10	0	0	1
9	10	11	1	0	1
10	9	10	0	2	1
11	10	9	10	9	11
9	10	11	9	99	11
10	9	9	11	10	10

O

X	X	X	X	X	X
X	10				X
X					X
X			9		X
X				10	X
X	X	X	X	X	X

10, 9, 11, 9, 99, 11, 11, 10, 10

sort

9, 9, 10, 10, 10, 11, 11, 11, 99

median

# Median Filter Properties

---


Non-linear

Does not spread the noise

Can remove spike noise

Expensive to run

# Median Filter



a b c

**FIGURE 3.37** (a) X-ray image of circuit board corrupted by salt-and-pepper noise. (b) Noise reduction with a  $3 \times 3$  averaging mask. (c) Noise reduction with a  $3 \times 3$  median filter. (Original image courtesy of Mr. Joseph E. Pascente, Lixi, Inc.)