

EECE 5639 Computer Vision I

Lecture 10

Hough Transform, RANSAC, Snakes

Next Class

Region segmentation

More Image Features

(Grouping edges)

Contours: Lines and Curves

Edge detectors find “**edgels**” (pixel level)

To perform image analysis :


edgels must be grouped into entities such as **contours** (higher level).

Canny does this to certain extent: the detector finds chains of edgels.

Line detection

Mathematical model of a line:

$$y = mx + n$$




$$y_1 = m x_1 + n$$

$$y_2 = m x_2 + n$$



$$y_N = m x_N + n$$

Image and Parameter Spaces



Line in Img. Space ~ Point in Param. Space

Looking at it backwards ...


Image space

Fix (m,n) , Vary (x,y) - Line

$$y = mx + n$$

Fix (x_1, y_1) , Vary (m,n) – Lines thru a Point

$$y_1 = m x_1 + n$$



Looking at it backwards ...

Parameter space


$$y_1 = m x_1 + n$$

Can be re-written as:

$$n = -x_1 m + y_1$$

Fix $(-x_1, y_1)$, Vary (m, n) - Line

$$n = -x_1 m + y_1$$



Img-Param Spaces

Image Space

Lines

Points

Collinear points

Parameter Space

Points

Lines

Intersecting lines

Hough Transform Technique

H.T. is a method for detecting straight lines (and curves) in images.



Main idea:

Map a difficult pattern problem into a simple peak detection problem

Hough Transform Technique

Given an edge point, there is an infinite number of lines passing through it (Vary m and n).

These lines can be represented as a line in parameter space.





Parameter Space

Hough Transform Technique

Given a set of collinear edge points, each of them have associated a line in parameter space.

These lines intersect at the point (m,n) corresponding to the parameters of the line in the image space.



Parameter Space

Hough Transform Technique

At each point of the (discrete) parameter space, count how many lines pass through it.

- Use an array of counters

- Can be thought as a “ **parameter image**”

The higher the count, the more edges are collinear in the image space.

- Find a peak in the counter array

- This is a “bright” point in the parameter image

- It can be found by thresholding

Practical Issues

The slope of the line is $-\infty < m < \infty$


The parameter space is INFINITE

The representation $y = mx + n$ does not express lines of the form $x = k$

Solution:

Use the “Normal” equation of a line:

$$y = mx + n$$



$$\rho = x \cos\theta + y \sin\theta$$

θ Is the line orientation

ρ Is the distance between the origin and the line

New Parameter Space

Use the parameter space (ρ, θ)

The new space is FINITE

$0 < \rho < D$, where D is the image diagonal.

$0 < \theta < \pi$

The new space can represent all lines

$y = k$ is represented with $\rho = k$, $\theta=90$

$x = k$ is represented with $\rho = k$, $\theta=0$

Consequence:

A Point in Image Space is now represented as a SINUSOID

$$\rho = x \cos\theta + y \sin\theta$$

Hough Transform Algorithm

Input is an edge image ($E(i,j)=1$ for edgels)

1. Discretize θ and ρ in increments of $d\theta$ and $d\rho$. Let $A(R,T)$ be an array of integer accumulators, initialized to 0.
2. For each pixel $E(i,j)=1$ and $h=1,2,\dots,T$ do
 1. $\rho = i \cos(h * d\theta) + j \sin(h * d\theta)$
 2. Find closest integer k corresponding to ρ
 3. Increment counter $A(h,k)$ by one
3. Find local maxima in $A(R,T)$

Hough Transform Speed Up

If we know the orientation of the edge – usually available from the edge detection step

We fix theta in the parameter space and increment **only one** counter!

We can allow for orientation uncertainty by incrementing a few counters around the “nominal” counter.

Hough Transform for Curves

The H.T. can be generalized to detect any curve that can be expressed in parametric form:

$$y = f(x, a_1, a_2, \dots, a_p)$$


a_1, a_2, \dots, a_p are the parameters

The parameter space is p -dimensional

The accumulating array is LARGE!

Generalizing the H.T.

The H.T. can be used even if the curve has not a simple analytic form!




$$x_c = x_i + r_i \cos(\alpha_i)$$

$$y_c = y_i + r_i \sin(\alpha_i)$$

1. Pick a reference point (x_c, y_c)
2. For $i = 1, \dots, n$:
 - a. Draw segment to P_i on the boundary.
 - b. Measure its length r_i , and its orientation α_i .
 - c. Write the coordinates of (x_c, y_c) as a function of r_i and α_i
 - d. Record the gradient orientation ϕ_i at P_i .
5. Build a table with the data, indexed by ϕ_i .

Generalizing the H.T.

Suppose, there were m **different** gradient orientations: ($m \leq n$)



$$x_c = x_i + r_i \cos(\alpha_i)$$


$$y_c = y_i + r_i \sin(\alpha_i)$$

ϕ_1	$(r^1_1, \alpha^1_1), (r^1_2, \alpha^1_2), \dots, (r^1_{n1}, \alpha^1_{n1})$
ϕ_2	$(r^2_1, \alpha^2_1), (r^2_2, \alpha^1_2), \dots, (r^2_{n2}, \alpha^1_{n2})$
.	.
.	.
.	.
ϕ_m	$(r^m_1, \alpha^m_1), (r^m_2, \alpha^m_2), \dots, (r^m_{nm}, \alpha^m_{nm})$

H.T. table

Generalized H.T. Algorithm:

Finds a rotated, scaled, and translated version of the curve:



$$x_c = x_i + r_i \cos(\alpha_i)$$

$$y_c = y_i + r_i \sin(\alpha_i)$$

1. Form an A accumulator array of possible reference points (x_c, y_c) , scaling factor S and Rotation angle θ .
2. For each edge (x, y) in the image:
 - a. Compute $\phi(x, y)$
 - b. For each (r, α) corresponding to $\phi(x, y)$ do:
 1. For each S and θ :
 - a. $x_c = x_i + r(\phi) S \cos[\alpha(\phi) + \theta]$
 - b. $y_c = y_i + r(\phi) S \sin[\alpha(\phi) + \theta]$
 - c. $A(x_c, y_c, S, \theta) ++$
3. Find maxima of A .

H.T. Summary

H.T. is a “voting” scheme

points vote for a set of parameters describing a line or curve.


The more votes for a particular set

the more evidence that the corresponding curve is present in the image.


Can detect MULTIPLE curves in one shot.

Computational cost increases with the number of parameters describing the curve.


Hough Transf. & Noise




tokens





votes





tokens



votes








Segmentation by Fitting

Fitting Lines

Using Least Squares Fitting Error:




$$\min \Phi = \sum_i (y_i - ax_i - b)^2$$

$$\frac{\partial \Phi}{\partial a} = 2a \sum_i x_i^2 - 2 \sum_i x_i (y_i - b) = 0$$


$$\frac{\partial \Phi}{\partial b} = -2 \sum_i (y_i - ax_i - b) = 0$$

$$\begin{aligned} \bar{x}\bar{y} &= \bar{x}^2 a + \bar{x}b \\ \bar{y} &= \bar{x}a + b \end{aligned}$$

$$\begin{bmatrix} \bar{x}\bar{y} \\ \bar{y} \end{bmatrix} = \begin{bmatrix} \bar{x}^2 & \bar{x} \\ \bar{x} & 1 \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix}$$



Line fitting can be max.
likelihood - but choice of
model is important




Fitting Lines

Using Total Least Squares Fitting Error:

$$\begin{aligned} \min \Phi &= \sum (ax_i + by_i + c)^2 \\ \text{s.t. } a^2 + b^2 &= 1 \end{aligned}$$

$$\min \Lambda = \sum_i (ax_i + by_i + c)^2 + \lambda(a^2 + b^2 - 1)$$



$$\frac{\partial \Lambda}{\partial a} = 0 \quad \frac{\partial \Lambda}{\partial b} = 0 \quad \frac{\partial \Lambda}{\partial c} = 0 \quad \frac{\partial \Lambda}{\partial \lambda} = 0$$

$$\begin{bmatrix} \bar{x}^2 - \bar{x}\bar{x} & \bar{x}\bar{y} - \bar{x}\bar{y} \\ \bar{x}\bar{y} - \bar{x}\bar{y} & \bar{y}^2 - \bar{y}\bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = \mu \begin{bmatrix} a \\ b \end{bmatrix}$$

Eigenvalue
Problem!

Who came from which line?

Assume we know how many lines there are - but which lines are they?

easy, if we know who came from which line

Possible strategies

Hough transform

Incremental line fitting

K-means

Algorithm 15.1: Incremental line fitting by walking along a curve, fitting a line to runs of pixels along the curve, and breaking the curve when the residual is too large

```
Put all points on curve list, in order along the curve
Empty the line point list
Empty the line list
Until there are too few points on the curve
    Transfer first few points on the curve to the line point list
    Fit line to line point list
    While fitted line is good enough
        Transfer the next point on the curve
            to the line point list and refit the line
    end
    Transfer last point(s) back to curve
    Refit line
    Attach line to line list
end
```

Algorithm 15.2: K-means line fitting by allocating points to the closest line and then refitting.

Hypothesize k lines (perhaps uniformly at random)

or

Hypothesize an assignment of lines to points
and then fit lines using this assignment

Until convergence

 Allocate each point to the closest line

 Refit lines

end

Robustness

As we have seen, squared error can be a source of bias in the presence of noise points

- One fix is EM - we'll not do this in this class

- Another is an M-estimator


 - Square nearby, threshold far away

- A third is RANSAC

 - Search for good points

Inliers-Outliers

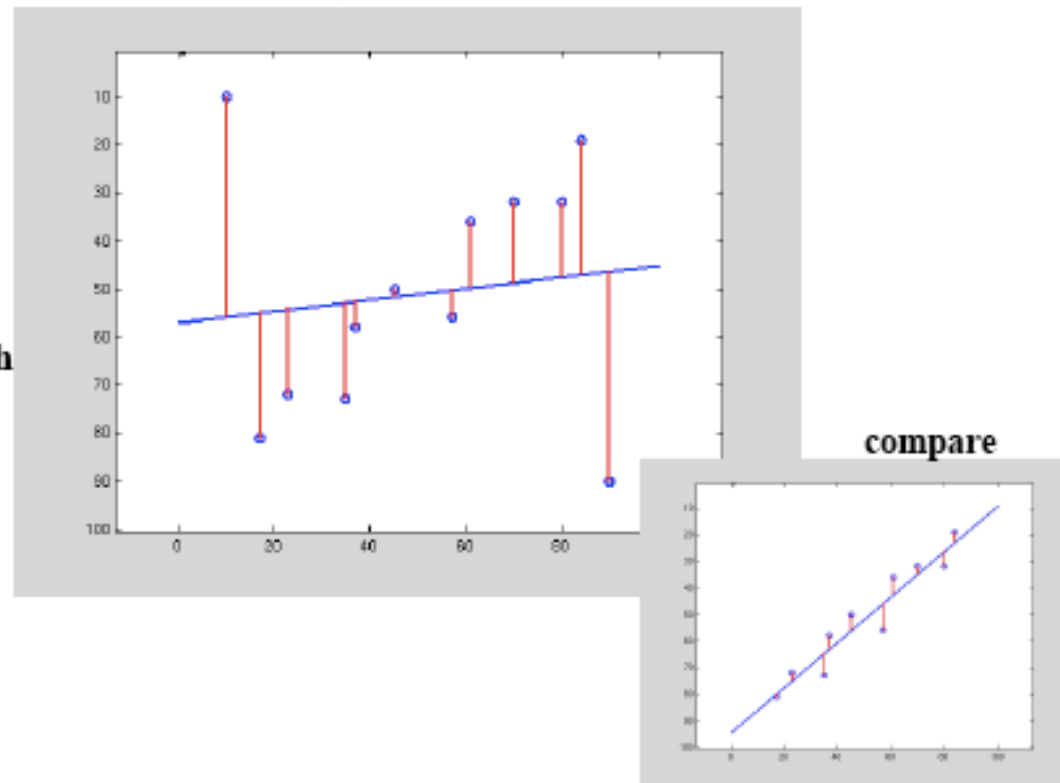
Loosely speaking, **outliers** are “bad data” points that do not fit the model. Points that fit the model are **inliers**.



Problems with Outliers

Least square estimation is very sensitive to outliers! :
Few outliers can GREATLY skew the result.


**Least squares
regression with
outliers**



Outliers are not the only problem

Multiple structures can also skew results.

The fitting procedure implicitly assumes ONE instance




Robust Estimation


Two steps:


Classify data into INLIERS and OUTLIERS


Use only INLIERS to fit the model

RANSAC is an example of this approach.









Robustness and M-estimators

LSE methods are very sensitive to outliers: one bad point can have tremendous effect on the solution.


An M-estimator is used to give different weights to the errors:

$$\min \sum_i \rho(r_i(x_i, \theta); \sigma)$$

**Residual error at
 x_i**

**Set of parameters
Being fitted**

**Parameter
Controlling
Error influence**




$$\rho(u; \sigma) = \frac{u^2}{u^2 + \sigma^2}$$


Issues:

1. Minimizing objective is not longer linear:
Solutions must be found interactively
2. Need to decide the scale parameter σ .


Too small sigma



Too large sigma



Good sigma value:




RANSAC

RANdom SAmple Consensus


RANSAC procedure (line example)




RANSAC procedure (line example)




RANSAC procedure (line example)




RANSAC procedure (line example)




RANSAC procedure (line example)




RANSAC procedure (line example)




RANSAC procedure (line example)




RANSAC procedure (line example)



RANSAC procedure (line example)



RANSAC procedure (line example)



RANSAC

Choose a small subset uniformly at random

Fit to that

Anything that is close to result is signal; all others are noise

Refit

Do this many times and choose the best

ISSUES

- How many times?
 - Often enough that we are likely to have a good line
- How big a subset?
 - Smallest possible
- What does close mean?
 - Depends on the problem
- What is a good line?
 - One where the number of nearby points is so big it is unlikely to be all outliers

Algorithm 15.4: RANSAC: fitting lines using random sample consensus

Determine:

- n — the smallest number of points required
- k — the number of iterations required
- t — the threshold used to identify a point that fits well
- d — the number of nearby points required
to assert a model fits well

Until k iterations have occurred

Draw a sample of n points from the data
uniformly and at random

Fit to that set of n points

For each data point outside the sample

Test the distance from the point to the line
against t ; if the distance from the point to the line
is less than t , the point is close

end

If there are d or more points close to the line
then there is a good fit. Refit the line using all
these points.

end

Use the best fit from this collection, using the
fitting error as a criterion

How Many Samples to Choose?

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of choosing one inlier

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of choosing s inliers

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of that one or more
points in the sample were
outliers (contaminated sample)

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of that N Samples
were contaminated

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of that AT LEAST
one sample of N Samples was
NOT contaminated

How Many Samples?

Choose N so that, with probability p , at least one random sample is free from outliers. E.g. $p = 0.99$

$$p = 1 - (1 - (1 - e)^s)^N$$

$$N = \frac{\ln(1 - p)}{\ln(1 - (1 - e)^s)}$$

proportion of outliers e							
s	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

Line example

12 pts: $n = 12$

Sample size: $s = 2$

Outliers 2: $e = 1/6 \rightarrow 20\%$

$N = 5$ gives 99% chance of getting a good sample
(trying every possible pair requires 66 trials!)


Acceptable Consensus Set

Typically, terminate when inlier ratio reaches expected ratio of inliers

$$T = (1 - e) \times \text{total number of data points}$$

After RANSAC

- RANSAC divides the data into inliers and outliers
- But it computes the estimate with the MINIMAL samples
- We can improve the result by estimating the model using all inliers:
- After RANSAC: estimate once more!



Fitting curves other than lines

In principle, an easy generalization

The probability of obtaining a point, given a curve, is given by a negative exponential of distance squared

- In practice, rather hard
 - It is generally difficult to compute the distance between a point and a curve

Active Contours





Deformable Contours

They are also called

Snakes

Active contours

Think of a snake as an elastic band:

of arbitrary shape

sensitive to image gradient

that can wiggle in the image

represented as a necklace of points

Active Contour Models

An important class of algorithms to find boundaries
Usually does not use prior knowledge of the shape
Poses the problem as an “optimization” problem

Introduction

How can we find the boundary of an object in an image?

One approach could be:

- Find edges

- Link the edges


Another possibility is to search for “smooth” boundaries:

- The boundary should “match” the image

- Can iteratively “improve”

Main Idea:

“Drop” a snake



Let the snake “wobble”, attracted by image gradient, until it glues itself against a contour

The Energy Functional

Associate to each possible shape and location of the snake a value E .

Values should be s.t. the image contour to be detected has the minimum value.

E is called the **energy** of the snake.

Keep wiggling the snake towards smaller values of E .

Energy Functional Design

We need a function that given a snake state, associates to it an Energy value E .

The function should be designed so that the snake moves towards the contour that we are seeking!

What moves the snake?

“Forces” applied to its points

Snake Energy

The total energy of the snake is defined as:

$$E_{total} = E_{internal} + E_{external}$$

The internal energy encourages smoothness

The external energy encourages closeness to edges

Forces moving the snake (External)

It needs to be attracted to contours:

Edge pixels must “pull” the snake points.

The stronger the edge, the stronger the pull.

The force is proportional to $|\nabla I|$

Forces preserving the snake (Internal)

The snake should not break apart!

Points on the snake must stay close to each other

Each point on the snake pulls its neighbors

The farther the neighbors, the stronger the force

The force is proportional to the distance $|P_i - P_{i-1}|$


Forces preserving the snake (Internal)

The snake should avoid “oscillations”

Penalize high curvature


Force proportional to snake curvature

Snake Forces



Snake “state”: Contour Parametrization

Consider a contour parametrization $c=c(s)$ where s is the “arc length”



Each point P_i on the contour
has coordinates $(x_i(s), y_i(s))$

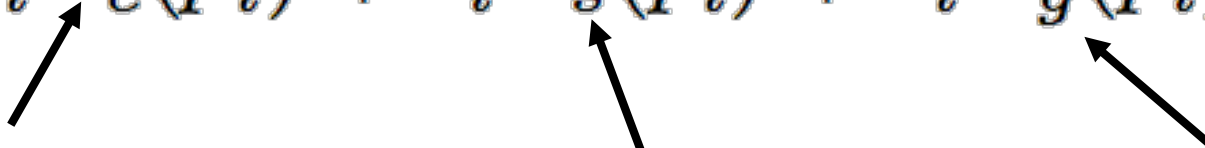
$$E = \int E_{int}(s) + E_{ext}(s) ds$$

Snake Energy Functional

Given a snake with N points p_1, p_2, \dots, p_N

$$E = \sum_{i=1}^N a_i E_c(p_i) + b_i E_s(p_i) + c_i E_g(p_i)$$

“Continuity” “Smoothness” “Edgeness”



a_i, b_i, c_i are “weights” to control influence

Continuity Term

Given a snake with N points p_1, p_2, \dots, p_N

Let d be the average distance between points

Distance between points should be kept close to average

Define the continuity term of the Energy Functional:

$$E_c(p_i) = (d - |p_i - p_{i-1}|)^2$$

$$p_i = [x_i \ y_i]$$

$$E_c = \left(d - \sqrt{(x_i - x_{i-1})^2 + (y_i - y_{i-1})^2} \right)^2$$

Smoothness Term

Given a snake with N points p_1, p_2, \dots, p_N

Curvature should be kept small

Define the smoothness term of the Energy Functional:

$$E_s(p_i) = \underbrace{|p_{i-1} - 2p_i + p_{i+1}|^2}_{\text{Second derivative}}$$

Second derivative

$$E_s = (x_{i-1} - 2x_i + x_{i+1})^2 + (y_{i-1} - 2y_i + y_{i+1})^2$$

Edgeness Term

Given a snake with N points p_1, p_2, \dots, p_N

Define the edgeness term of the Energy Functional:

$$E_g(p_i) = -|\nabla I(p_i)|$$




$$\nabla I(p_i) = [G_x(p_i) \ G_y(p_i)]$$

$$|\nabla I(p_i)| = \sqrt{G_x(p_i)^2 + G_y(p_i)^2}$$

Magnitude of the gradient should be LARGE


Relative Weighting

The weights control the smoothness and stiffness of the snake



Greedy Algorithm

Each point moves within a small window to minimize the energy



Compute the new energy for each candidate location
Move the point to the one with the **minimum** value

Keeping corners ...

Before starting a new iteration:

Search for “corners”:

- max curvature

- large gradient

Corner points should not contribute to the energy (set $b_i = 0$)

Implementation Considerations

To avoid numerical problems, the terms of the energy function should be normalized.

E_c and E_s are normalized by their maximum in the neighborhood

E_g is normalized as $|\nabla| - m| / (M - m)$

M and m are the max and min value of the gradient magnitude in the neighborhood

Snake Algorithm

Input:

gray scale image I

a chain of points p_1, p_2, \dots, p_N

f is the fraction of points that must move to start a new iteration

$U(p)$ is a neighborhood around p

d is the average distance between snake points.

Snake Algorithm

1. While the fraction of moved points $> f$
 1. For $i=1,2,\dots,N$
 1. find a point in $U(p_i)$ s.t. the energy is minimum,
 2. move p_i to this location
 2. For $i=1,2,\dots,N$
 1. Estimate the curvature $k=|p_{i-1}-2p_i+p_{i+1}|$
 2. Look for local max, and set $b_{\max} = 0$
3. Update d

Problems with Snakes

Smoothness does not always capture all prior knowledge

User must define the weights

Snakes might oversmooth boundaries

Not trivial to prevent curve self intersecting

